

Week 8: Logistic Regression

10/17/2019

Jake Campbell


Moving On From Linear Problems

- Up to now, we've only looked at modeling continuous response variables
 - Miles per gallon, rate of crime, etc.
- There are several problems where we need to model a categorical response
 - Does a team win a game?; Does a student accept an offer of admission?
- These problems involving categorical response variables are classification problems

Issues With Modeling This Problem Linearly

- Let's say we were trying to model a yes/no problem
- We can treat all yes as **1** and no as **0** and predict the probability of either a **1** or **0**
- A simple linear line isn't really flexible to this scenario
- Predictions can go above and beyond **0** or **100%** probability

Issues With Modeling This Problem Linearly



Logistic Regression

- We can deal with binary classification using the logistic function
 - Can take any input and return a value between 0 and 1

$$P = \frac{e^{(B_0 + B_x)}}{1 + e^{(B_0 + B_x)}}$$

- By using some simple math, we can edit the logistic function to get something similar to the linear regression equation
 - Instead of predicting y, we are predicting the log odds of y

$$\log\left(\frac{P}{1 - P}\right) = B_0 + B_x$$

Odds and Probability

- We aren't predicting probability in logistic regression, we are predicting the log of the odds ratio
- Probability can be easily changed to odds and vice versa
- The odds ratio is simply the odds of event A happening in the numerator and the odds of it not happening in the denominator
- For example, if a football team had an 80% chance of winning, their odds of winning would be $.8 / 1 - .8$ or $4 / 1$
 - The team has 4 to 1 odds of winning; if the game was played 5 times, the team would win 4 times

Logistic Regression in R

- We will use `glm()` or generalized linear model to fit logistic regression models
- `glm` can fit several non-linear functions
- Specify the argument `family = "binomial"` to perform logistic regression
- `summary()` can be used similarly to how we use it for `lm()` models

Logistic Regression Output

- Like in our `lm` output, we get coefficient estimates and p-values.
- Coefficient estimates are the increase in the log odds of y for a one unit increase in x
- Deviance is a measure of model fit
 - The lower the better
- Residual deviance is the model's deviance, while null deviance is the deviance with only an intercept
 - Similar to F-test in `lm` output

Interpreting Logistic Regression Coefficients

- Remember that we are predicting the log odds of event **y** occurring
 - A one unit increase in **x**, leads to an increase in the log odds of event **y** occurring by its coefficient estimate
- We can take the exponential of our coefficient estimates to convert them to odds

```
exp(coef(tumor.glm))
```

```
## (Intercept)          V1          V3          V4          V5
## 0.0000411763 1.7067054139 1.3745942848 1.3908643372 1.1008866850
##          V6          V7          V8          V9
## 1.4667887059 1.5628015794 1.2368248052 1.7058506868
```

- For example, a one-unit increase in **V1** would up the odds of the tumor being malignant by about **70%**

Logistic Regression Assumptions

- Response is binary (two classes)
- Independent observations
- No multicollinearity issues
- Large sample sizes
 - Logistic regression fit by maximum likelihood which requires larger sample sizes

The Role of Polynomials

- We don't expect variables to have a linear relation with the response; we do expect variables to have a linear relation with the log odds
- The log odds of a model can be plotted against an independent variable
- Relationships that don't appear linear might be improved with the use of polynomials

Measuring Accuracy

- We can make class probability predictions using `predict()` and specifying `type = "response"`
- Without specifying type, we get log odds predictions
- We can round the predicted probabilities to the nearest whole number and measure how accurate our model is
- Straightforward method of determining how good our model is