

Week 10: Model Validation

03/31/2020

Jake Campbell

The Importance of Cross-validation

- Building a model is great, but what we really care about are the predictions
- Gauging model performance on predictions on data the model was trained on is a no-no
 - Your model has already seen this data; predictions are likely to be overly-optimistic
- A real test of model performance is looking at how well the model predicts data it hasn't seen

Training/Testing Approach

- A simple method of cross-validation is to hold out some data from the model training process
- This will act as a test set
- Generally speaking, you would want the majority of observations to be in your training set (70-80%)
- `createDataPartition` can be used to create a stratified random sample of indices to use as a training set

```
data_split <- createDataPartition(Sonar$Class, p = .75, list = F)
# Index Sonar by data_split to get a training and testing set
training <- Sonar[data_split, ]
testing <- Sonar[-data_split, ]
```

Issues with the Training/Testing Approach

- Depending on which rows are randomly chosen in our dataset, we can have varied results
 - Maybe the model was trained on the easy to predict observations
- We are leaving out data we could be training the model on
- Can be easily influenced by researcher bias
 - If we have just one testing set, we can continue to adjust model parameters to try and make our test set look ideal
 - If we don't like the results we get for our test set, we can just resample post-hoc until we like our test set results

Other Cross-validation Methods

- Leave One Out
 - A model is built on all observations except one; this one observation is used as the test set
 - This process is repeated for each row and the error is averaged across all observations
- Bootstrap
 - The dataset is continuously sampled with replacement creating several training sets
 - Test sets are the unused observations from each sample
- K-fold Cross Validation
 - The dataset is split into k-folds with a model being trained on each combination of folds, leaving one fold out to act as a testing set each time
 - Results are averaged across each of the test sets

caret

- Classification And REgression Training
- Allows us to perform other validation methods using several different types of models
- Full list of models available and what package they are from available on caret's documentation
 - <https://topepo.github.io/caret/available-models.html>
- Note that the model created using caret is a caret object and not a specific model's object
- Also has several preprocessing functions
 - Splitting, scaling, centering, etc.

Creating Models with **caret**

- Model syntax is the same as other models, but there are additional arguments to be made
- Input model formula into caret's **train()** function
- Specify **method** argument to specify what type of model you want caret to use
- Specify **trControl** using the **trainControl()** function
 - We can input the cross validation method within **trainControl** function
- Specify a grid of tuning parameters using the **tuneGrid** argument
 - Must be input as a dataframe where each parameter is a column

Common Cross-validation Methods with **caret**


- We need to specify the cross validation method within the **trainControl()** function
 - Specify with **method** argument
- K-Fold CV: **method = "repeatedcv"**
 - Specify number of folds and number of repeats
- Leave One Out: **method = "LOOCV"**
- Bootstrap: **method = "boot"**
 - Specify number of resamples

caret Output

- The final model created by caret is a model using all data observations
- By specifying `savePredictions = T` in the `trainControl` function, we save results from each cv fold
 - We can use this to identify an expected distribution of what error metric to expect

caret Output

```
ggplot(data = sonar_glm_cv$resample, aes(x = Accuracy)) +  
  geom_density(alpha = .2, fill="red")
```



caret and confusionMatrix

- `caret` has a `confusionMatrix` function
 - Creates a confusion matrix as well as gives several different accuracy measurements
 - Specify `data` as your predictions and `reference` as the actual values
 - Set the positive class with the `positive` argument

caret and confusionMatrix

```
confusionMatrix(data = class_predictions, reference = testing$Class,  
                 positive = "R")
```

```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction  M   R  
##           M 18   5  
##           R   9 19  
##  
##           Accuracy : 0.7255  
##           95% CI : (0.5826, 0.8411)  
## No Information Rate : 0.5294  
## P-Value [Acc > NIR] : 0.003347  
##  
##           Kappa : 0.4541  
##  
## Mcnemar's Test P-Value : 0.422678  
##  
##           Sensitivity : 0.7917  
##           Specificity : 0.6667  
##           Pos Pred Value : 0.6786  
##           Neg Pred Value : 0.7826
```

caret for Future Use

- In addition to being a tool for cross validation, **caret** has importance in model selection
- Not really pertinent to this class, because models we go over have no additional tuning parameters
- Several models can be built based off of different tuning parameters
 - Ex: Boosted trees can be built with a different number of tree based models; caret can build models at different intervals for number of trees (**100**, **500**, **1000**, etc.) and the models can be compared based on different metrics